

User *Manual*

Assistive AI feedback for actors' headshot portraits.
Structured evaluation across technical, aesthetic, and casting
dimensions — for photographers, actors, and casting
professionals.

Technical Quality

Aesthetic & Market Fit

Casting Attributes

Assistive Only

CONTENTS

Table of Contents

Everything you need to evaluate, curate, and export headshots with confidence.

01	Introduction & Positioning	3
02	The Three-Dimension System	4
03	Getting Started	5
04	Understanding the Evaluation	6
05	Market Contexts	7
06	Shortlist & Curation	8
07	Comparison View	9
08	Keyboard Shortcuts	10
09	Mobile Experience	10
10	Ethics & Fairness	11
11	Technical Architecture	12
12	API Reference	13
13	Deployment	14

HOW TO READ THIS MANUAL

Sections 1–5 are for all users. Sections 6–9 cover curation and power features. Sections 10–13 are for developers and operators deploying or extending the system.

SECTION 01

Introduction & Positioning

An assistive tool that evaluates a photograph — never a person.

Spotlight Review is an AI-powered headshot evaluator designed for photographers, actors, and casting professionals. It analyses casting headshots across three structured dimensions — technical quality, aesthetic and market fit, and casting attributes — and returns actionable, transparent feedback with explicit confidence bands.

The system is built on a fundamental principle: it evaluates a **photograph**, never a **person**. Every output is phrased as a property of the image and the casting context. The tool never judges attractiveness as a headline metric, never issues a verdict on someone's worth, and always communicates uncertainty. It exists to support human judgement, not to replace it.

"Assistive feedback on photographs, for humans who remain in charge."

— Product positioning principle

Who it's for



Actors & Agents

Curate a portfolio of 50+ portraits down to a top-8 shortlist optimised for range and market fit.



Photographers

Cull 200+ frames after a session. Flag technical issues, rank by quality, and coach with specific notes.



Casting Directors

Surface the strongest headshots in a submission batch with structured feedback and type labels.



Portfolio Managers

Export a ranked, annotated shortlist as a print-ready PDF for client review or platform submission.

WHAT THIS TOOL IS NOT

This is **not** "AI that tells you how attractive you are." It does not assess a person's beauty. It does not serve as an automated screening gate. It does not make casting decisions. Final judgement always rests with the human professional.

SECTION 02

The Three-Dimension System

Every headshot is assessed across three independent dimensions, each broken into sub-scores.

Rather than a single opaque score, Spotlight Review breaks each evaluation into three dimensions that map to the real concerns of casting and photography. Each dimension produces sub-scores (0–1), issue flags, and contributes to a weighted composite overall score (0–100). The weights are configurable per market context.

1

Technical Quality

Is the photograph well-made? The craft of the image — exposure, focus, noise, color, and dynamic range.

Exposure

Eye Focus

Low Noise

Color Balance

Dynamic Range

Eye focus is the single most important technical signal in a headshot — the near eye should be the sharpest point in the frame. Flags include: `soft_focus`, `backlit`, `noisy`, `underexposed`, `overexposed`, `color_cast`, `clipped_highlights`, `motion_blur`.

2

Aesthetic & Market Fit

Does the photograph suit its intended market? Composition, background, expression, wardrobe, and believability.

Composition

Background

Expression Fit

Wardrobe Fit

Believability

Expression and wardrobe fit are judged **relative to the selected market context** — the same expression that scores high in "Drama" may score lower in "US Commercial."

Flags include: `busy_background`, `tight_crop`, `loose_crop`, `distracting_wardrobe`, `flat_lighting`, `overstylised`.

3

Casting Attributes

Who could this headshot play? Apparent age, gender presentation, type labels, and expression readability.

Age Range

Gender Presentation

Type Labels

Expression Readability

Type labels include: `commercial`, `romantic_lead`, `villain`, `best_friend`, `character_actor`, `corporate`, `authority`, `dramatic_lead`. Gender presentation includes an explicit `uncertain` class with a confidence score — the model never forces a binary when it's genuinely unsure.

SECTION 03

Getting Started

From upload to evaluation in under a minute.

1

Set the market context

Choose one of five market lenses before uploading. The same headshot scores differently per market — this is legitimate, not a bug. Pick the lens you're curating for: US Commercial, UK Theatrical, Drama, Comedy, or Film & TV.

2

Upload your headshots

Drag and drop a batch of portraits (JPG, PNG, WebP) onto the upload zone, or click "Browse files." Images are automatically resized to max 1024px before upload to keep payloads small. Alternatively, click "Try sample headshots" to evaluate 6 pre-loaded portraits instantly.

3

Wait for evaluation

Evaluations process sequentially (one at a time, ~12 seconds each) to respect API rate limits. A sticky progress bar shows processed count, percentage, and ETA. You can pause, resume, or cancel at any time.

4

Review and curate

Each portrait card shows the overall score, top type label, apparent age, and any flags. Click any card to open the full detail dialog with radar chart, per-dimension breakdown, casting attributes, and structured feedback. Pin, reject, override, or compare as you go.

5

Export your shortlist

Click "Export PDF" to generate a print-ready document of your top-8 shortlist with scores, flags, and notes. The PDF opens in a new window — use your browser's "Save as PDF" to download.

PERSISTENCE

Your portraits, evaluations, pins, and overrides survive a page refresh. Uploaded images are stored in the browser's IndexedDB; metadata and evaluations are stored in localStorage. No data is lost if you accidentally close the tab.

SECTION 04

Understanding the Evaluation

Every score is transparent, explainable, and accompanied by a confidence band.

The overall score

The 0–100 composite score is computed from the three dimensions using per-market weights. Technical quality typically carries 25–35% weight, aesthetic and market fit 35–45%, and casting fit 20–40%. A small flag penalty (2 points per technical flag, capped at 6) is subtracted. The exact breakdown is shown in the detail dialog under "How this score was computed."



Score bands: **80–100 Excellent** · **65–79 Strong** · **50–64 Competent** · **0–49 Needs work**

Confidence and uncertainty

Every evaluation ships with a confidence band (0–1). When confidence drops below 50%, the UI displays an amber notice recommending a human reviewer for the final call. This is a core fairness feature — the model is honest about its own reliability rather than presenting every score as equally trustworthy.

Structured observations vs. model notes

The Feedback tab in the detail dialog shows two types of feedback:

Type	Characteristics
Structured observations	Deterministic, auditable feedback generated from the structured scores and flags. Each flag maps to a fixed, human-reviewed sentence. Impossible to hallucinate. Examples: "Eyes are tack-sharp," "Background is busy — consider a cleaner backdrop."
Model notes	Free-form observations from the vision model. Supplementary to the structured observations. More nuanced but less auditable. Shown separately so users always have a trustworthy baseline.

The radar chart

The detail dialog's Overview tab includes a 10-axis radar chart that visualises all sub-scores simultaneously — five technical axes and five aesthetic axes. This lets you spot patterns at a glance: a portrait with high exposure but low eye focus, or high composition but low background cleanliness, tells a specific story about what to fix.

SECTION 05

Market Contexts

The same headshot reads differently per market. Pick the lens you're curating for.

Spotlight Review supports five market contexts. Each context adjusts the evaluation weights, the expression-fit and wardrobe-fit criteria, and the composite score. Switching market is instant if a cached evaluation exists for that market — otherwise the system re-evaluates with the new lens.

US Commercial

Bright, friendly, approachable. Even lighting, clean backdrop. Rewards warm expressions with visible teeth and engaged eyes. Penalises heavy shadow and overly intense expressions.

UK Theatrical

Tolerates shadow, intensity, and character. Mid-to-high contrast is acceptable. Neutral or thoughtful expressions are valued. Penalises overly "smiley commercial" styling and flat snapshot lighting.

Drama

Rewards emotional depth, strong readable eyes, deliberate contrast, and a sense of interior life. Wardrobe may be muted and characterful. Penalises stiff "say cheese" smiles and flat even lighting.

Comedy

Rewards warmth, energy, and an alive, expressive face. A genuine engaged smile and connecting eyes are key. Penalises deadpan neutral that reads cold and harsh unflattering shadow.

Film & TV

Rewards naturalistic believability — mid contrast, real skin texture, subtle expression, wardrobe that could exist in the world of the story. Penalises overt "modelling shoot" stylisation, heavy retouching cues, and overly graphic backdrops.

MARKET SWITCHING & CACHING

When you change market, any portrait that already has a cached evaluation for that market is updated instantly — no re-evaluation needed. Only portraits without a cached result for the new market are sent to the API. This makes market comparison a genuine exploration tool rather than a waiting game.

SECTION 06

Shortlist & Curation

Pin, reject, override, and let the smart shortlist balance score with range.

The smart shortlist

The suggested shortlist ranks your top-8 portraits, but it doesn't simply pick the 8 highest scores. It balances score with **expression and type diversity** — a mix of commercial, dramatic, and character types at slightly lower scores beats eight near-identical frames. Pinned portraits always win: if you pin a portrait, it's guaranteed a spot in the shortlist regardless of score.

Pin, reject, and override

Action	Shortcut	Effect
Pin	F	Force-includes the portrait in the shortlist. Pinned items always rank, regardless of score. Unpins if already pinned.

Reject	X	Removes the portrait from the shortlist and greys it out in the grid. Toggle "show rejected" in the toolbar to review. Un-rejects if already rejected.
Override score	—	Drag the slider in the detail dialog to set a custom score. Marked as "score overridden" on the card. Your override replaces the model's score in the shortlist ranking.
Remove	—	Deletes the portrait entirely. An undo toast appears for 6 seconds — click "Undo" to restore.

Multi-select mode

Click "Select" in the toolbar to enter multi-select mode. Click portrait cards to select them (a checkmark appears), then use the bulk action bar to pin, reject, or remove all selected at once. "All" selects every non-rejected portrait; "None" clears the selection.

Outcome tracking

In the detail dialog, you can optionally mark a headshot's submission outcome: **Booked**, **Audition**, or **No response**. This is a secondary, noisy training signal — it helps correlate headshot qualities with real-world results over time. Outcomes appear as a badge on the portrait card and are logged (with consent) to the feedback database.

YOUR EDITS TEACH THE TOOL

Every pin, reject, and override is framed as "teaching the tool." If you've granted consent, these edits are logged anonymously (image hash + action only, never the image itself or your identity) as training signal to improve the evaluator over time. You can revoke consent at any time.

SECTION 07

Comparison View

Side-by-side analysis of up to four portraits, with the system's reasoning visible.

The comparison view is where the system earns trust. Add up to four portraits to the compare set using the compare button on any portrait card (the git-compare icon). Then click "Compare" in the toolbar to open the side-by-side dialog.

What you see

- **Side-by-side images** with score rings and quality-band labels
- The **"Best" badge** — automatically highlights the highest-scoring portrait in the set
- **Overlaid radar charts** — each portrait's 10-axis radar is shown so you can visually compare dimension profiles
- **Key metrics** — eye focus, background, expression fit, and believability as horizontal bars for direct comparison
- **Top type label** and flags for each portrait
- **"View full details"** button to jump to any portrait's detail dialog

WHY COMPARISON MATTERS

For a curation tool, comparison is the core interaction — not a nice-to-have. When a photographer has 200 frames from a session, they don't just want individual scores; they want to know *why one ranks higher than another*. The comparison view makes the model's reasoning visible: you can see exactly which dimensions differ, which flags fire on one but not the other, and whether the ranking aligns with your professional judgement.

Compare limits

You can compare up to **4 portraits** at a time. The compare button on cards is disabled once you reach 4. Remove a portrait from the compare set (either from the card or inside the comparison view) to add a different one.

SECTION 08

Keyboard Shortcuts

Power-user culling at the speed of thought.

Spotlight Review supports keyboard shortcuts for fast, keyboard-driven culling — essential when working through 200+ frames. Shortcuts are disabled while typing in an input field or when a dialog is open (except Escape, which closes the active dialog).

Key	Action
 	Select the next portrait (wraps around)

K ↑	Select the previous portrait (wraps around)
Enter	Open the detail dialog for the selected portrait
F	Toggle pin (flag) on the selected portrait
X	Toggle reject on the selected portrait
R	Retry evaluation on the selected portrait (useful if it errored)
?	Toggle the keyboard shortcuts help dialog
Esc	Close the active dialog or detail view

WORKFLOW TIP

For a fast cull: press J to move to the next portrait, F to flag keepers, X to reject duds, and Enter to inspect any borderline frame in detail. You can process dozens of portraits per minute without touching the mouse.

SECTION 09

Mobile Experience

Designed for actors curating on the go.

Spotlight Review is fully responsive and optimised for mobile use. The layout adapts at each breakpoint: the portrait grid collapses from 5 columns to 2, the market selector and upload zone stack vertically, and the header collapses to essentials.

Mobile-specific features

- **Floating shortlist button** — a floating "Shortlist N" button sits in the bottom-right corner, opening a bottom-sheet drawer with your ranked top-8. This solves the problem of the shortlist sitting below the entire grid on mobile.
- **44px tap targets** — all card action buttons (pin, reject, compare, remove) meet the Apple/Google 44px minimum touch-target guideline.
- **Tabbed detail dialog** — on mobile, the detail dialog uses a 5-tab layout (Overview, Technical, Aesthetic, Casting, Notes) instead of a long scroll, keeping each section digestible.

- **Camera upload** — the "Browse files" button opens the native mobile file picker, which on iOS includes the camera. You can shoot and evaluate directly.

SECTION 10

Ethics & Fairness

Subjectivity honoured, bias under watch, human always in charge.

Ethics is not a section bolted on at the end — it shapes the architecture. Spotlight Review is designed around three ethical pillars that govern every design and engineering decision.

PILLAR 1 — SUBJECTIVITY, HONOURED

Beauty and casting suitability are not universal. Different markets and genres have different standards. The system is configurable per market context, surfaces confidence on every output, and explicitly recommends human review when the model is outside its reliable range. When the model is unsure, it says so — it never presents a low-confidence score as equally trustworthy.

PILLAR 2 — BIAS UNDER WATCH

Models trained on skewed data can under-score certain ethnicities, ages, body types, or gender presentations. The design calls for: diverse training data collection, bias testing and monitoring by demographic group on a frozen benchmark, and surfacing per-subgroup confidence and reliability flags. The gender presentation output includes an explicit "uncertain" class — the model never forces a binary when it's genuinely unsure, and the UX prompts users to self-identify.

PILLAR 3 — ASSISTIVE, NOT A VERDICT

This tool evaluates a photograph, never a person. It exists to support photographers, actors, and casting — never to replace human judgement or act as an automated screening gate. The product SLA is advisory. The system never outputs an "attractiveness" score as a headline. Override and pin are first-class features, not workarounds. The framing throughout is "assistive feedback for photographers and actors," not "AI that tells you how attractive you are."

Consent and data

Before any user edits are logged as training signal, the consent modal appears. Users choose to "help improve the tool" or "keep edits private." Only the image hash and the action (pin, reject, override) are stored — never the image itself, never the user's identity. Consent can be revoked at any time. All logging is gated by the consent record in the database.

Guardrails


- The system **never** outputs an attractiveness score as a headline
- The system **never** auto-rejects a portfolio — it ranks and annotates
- Age and gender outputs are presented as ranges with uncertainty
- A **public feedback channel** lets users flag scores they believe are wrong
- The internal review board signs off on each model release

SECTION 11

Technical Architecture

A multi-provider vision system with a durable structured-output contract.

System flow



JSON extract + coerce + composite score



HeadshotEvaluation JSON response

Provider abstraction

The vision backend is swappable via the `VISION_PROVIDER` environment variable. Two providers are implemented:

Provider	Model	Notes
<code>gemini</code> (default)	Google Gemini 2.5 Flash	Free tier (1,500 req/day). Public API, reachable from Vercel. Uses native JSON response schema for reliable structured output.
<code>zai</code>	GLM-4V (ZAI SDK)	Internal Z.ai API, only reachable from inside Z.ai's network. Used in the sandbox environment. Includes 429 retry-with-backoff.

Structured output contract

The `HeadshotEvaluation` type is the durable asset — the backend model is swappable, but the output schema stays constant. This means the frontend never changes when switching providers. The contract includes: `overall_score` (0–100), `overall_confidence` (0–1), three dimension blocks with sub-scores and flags, casting attributes, narrative feedback array, market context, model version, and disclaimer.

Client architecture

The frontend uses a Zustand store with persistence (localStorage for metadata, IndexedDB for image blobs). The store manages: the portrait collection, evaluation queue (sequential with pause/resume/cancel), per-market client-side cache, multi-select state, compare set, consent gating, undo stack, and outcome tracking. The evaluation queue processes one image at a time with a 1.2s gap to respect API rate limits.

Database

Prisma ORM with SQLite. Four models: `EvaluationCache` (image hash + market → cached result), `FeedbackLog` (user edits as training signal), `Outcome` (booked/audition/no-response), `ConsentRecord` (per-session consent). On Vercel, SQLite is ephemeral (serverless filesystem) — the app gracefully handles DB failures, and client-side caching ensures the UX is unaffected.

SECTION 12

API Reference

Five endpoints powering the evaluation, feedback, consent, and export flows.

POST /api/evaluate

Evaluates a headshot image. Accepts a data URL or remote URL and a market context. Returns the full `HeadshotEvaluation` object with composite score, sub-scores, flags, casting attributes, and narrative feedback. Results are cached server-side by image hash + market.

Request: `{ "image": "data:image/jpeg;base64,...", "market": "us_commercial", "imageHash": "..."}`

Response: `{ "evaluation": HeadshotEvaluation, "imageHash": "..."}`

POST /api/feedback

Logs a user edit (pin, reject, override, remove, compare) as training signal. Only persists if the session has granted consent. Always returns 200 so the UX never blocks on logging.

Request: `{ "sessionKey": "...", "imageHash": "...", "market": "...", "action": "pin", "value": "optional"}`

POST /api/outcome

Records a submission outcome (booked, audition, no_response) for a headshot. Secondary, noisy training signal. Requires consent.

Request: `{ "sessionKey": "...", "imageHash": "...", "market": "...", "result": "booked", "note": "optional"}`

GET / POST /api/consent

Checks or sets consent for a session. `GET` returns the current consent state; `POST` grants or revokes. The session key is an anonymous client-generated ID.

GET: `/api/consent?sessionKey=...` → `{ "ok": true, "granted": true }`

POST: `{ "sessionKey": "...", "granted": true }`

POST /api/export-shortlist

Generates a standalone, print-optimised HTML document for the shortlist. The client opens it in a new window and calls `print()` — the user can then "Save as PDF." The document includes ranked cards with images, scores, flags, type labels, and narrative notes.

Request: `{ "items": [...], "market": "us_commercial", "sessionName": "Shortlist 2024-01-15" }`

Response: HTML document (Content-Type: text/html)

SECTION 13

Deployment

Running Spotlight Review on Vercel with Google Gemini's free tier.

Live deployment

The production deployment runs on Vercel serverless (US region) using Google Gemini 2.5 Flash as the vision provider. The Gemini free tier allows 1,500 evaluations per day at no cost — sufficient for a demo or small-scale production use.

Component	Value
Production URL	<code>spotlight-review.vercel.app</code>
Source	<code>github.com/Optitude-AI/spotlight-review</code>
Vision provider	Google Gemini 2.5 Flash (free tier)
Framework	Next.js 16 (App Router, Turbopack)
Hosting	Vercel serverless functions
Database	SQLite (ephemeral on serverless)

Environment variables

Variable	Required	Purpose
<code>VISION_PROVIDER</code>	Yes	<code>gemini</code> or <code>zai</code> . Default: <code>gemini</code>
<code>GEMINI_API_KEY</code>	If <code>gemini</code>	Free key from <code>aistudio.google.com/apikey</code>
<code>DATABASE_URL</code>	Yes	SQLite path. Use <code>file:/tmp/custom.db</code> on Vercel
<code>ZAI_*</code>	If <code>zai</code>	ZAI credentials (<code>baseUrl</code> , <code>apiKey</code> , <code>token</code> , <code>chatId</code> , <code>userId</code>)

Deploy script

A non-interactive deploy script is included at `deploy-to-vercel.sh`. It sets all environment variables and deploys to production in one command:

```
VERCEL_TOKEN=<token> GEMINI_API_KEY=<key> bash deploy-to-vercel.sh
```

Local development

For local development, the ZAI provider is used by default in the Z.ai sandbox (where `/etc/.z-ai-config` exists). To use Gemini locally, set `VISION_PROVIDER=gemini` and `GEMINI_API_KEY` in `.env`. Run `bun run dev` to start the dev server on port 3000.

SQLITE ON VERCEL

SQLite is ephemeral on Vercel's serverless filesystem — server-side caching and feedback logging won't persist across function invocations. The app gracefully handles this: client-side caching (in the browser) still works, and all API routes catch DB failures silently. For persistent server-side storage, consider Vercel Postgres or Turso (libSQL).

GETTING HELP

Source code and issues: github.com/Optitude-AI/spotlight-review · Live app: spotlight-review.vercel.app